

COMPANY_ID	UNIT	UNIT_COST
10	12	1.15
12	12	1.05
14	18	1.31
18	18	1.34
11	24	1.15
16	12	1.31
10	12	1.15
12	24	1.3
18	6	1.34
18	12	1.35
14	12	1.95
21	18	1.36

If we run the query below:

```
SELECT
  company_id,
  units,
  unit_cost
FROM
  costing
WHERE
  company_id = 18
```

The database would have to go through all the records one by one, in the order, they appear to search for the matching datasets. It will become more and more time-consuming once the table size increases.

COMPANY_ID	UNIT	UNIT_COST
10	12	1.15
10	12	1.15
11	24	1.15
11	24	1.15
12	12	1.05
12	24	1.3
12	12	1.05
14	18	1.31
14	12	1.95
14	24	1.05
16	12	1.31
18	18	1.34
18	6	1.34

So here, ***indexing comes into the picture.***

How indexing actually works?

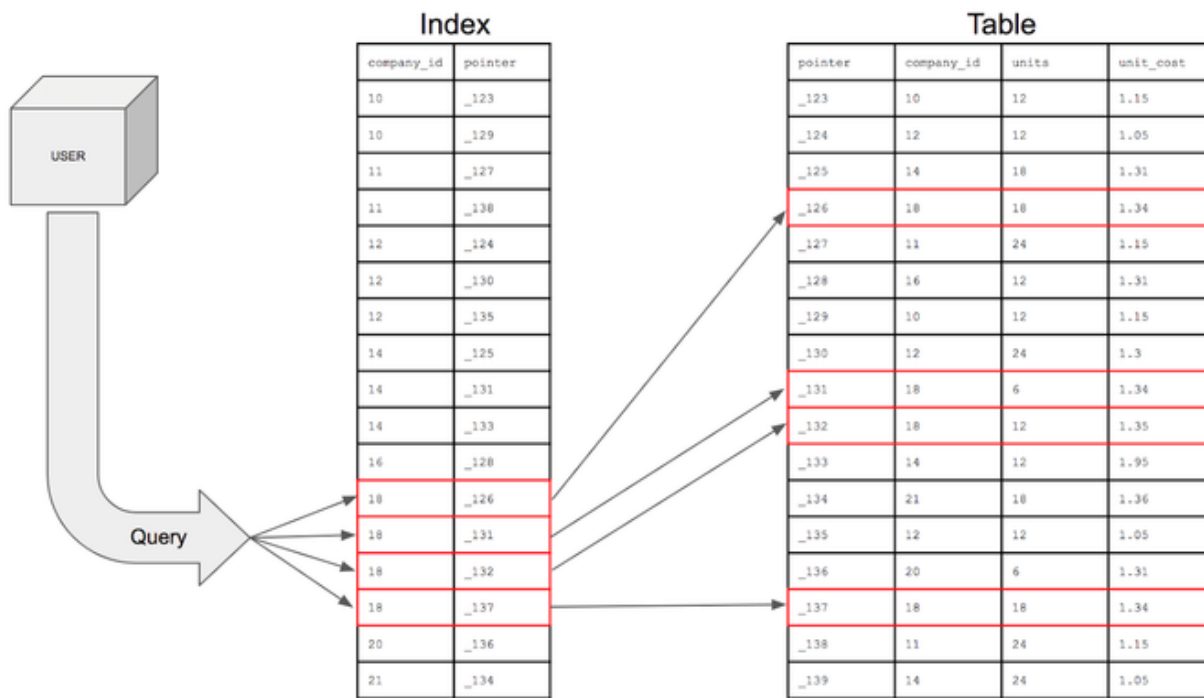
An index is a data structure, (most commonly its B-tree Its balanced tree, not binary tree) that stores the value for a specific column in a table.

The database table does not reorder itself when we put an index on a column to optimize the query performance.

So the index would look like this -

COMPANY_ID	POINTER
10	_123
10	_129
11	_127
11	_138
12	_124
12	_130
12	_135
14	_125
14	_131
14	_133
16	_128
18	_126
18	_131
18	_132

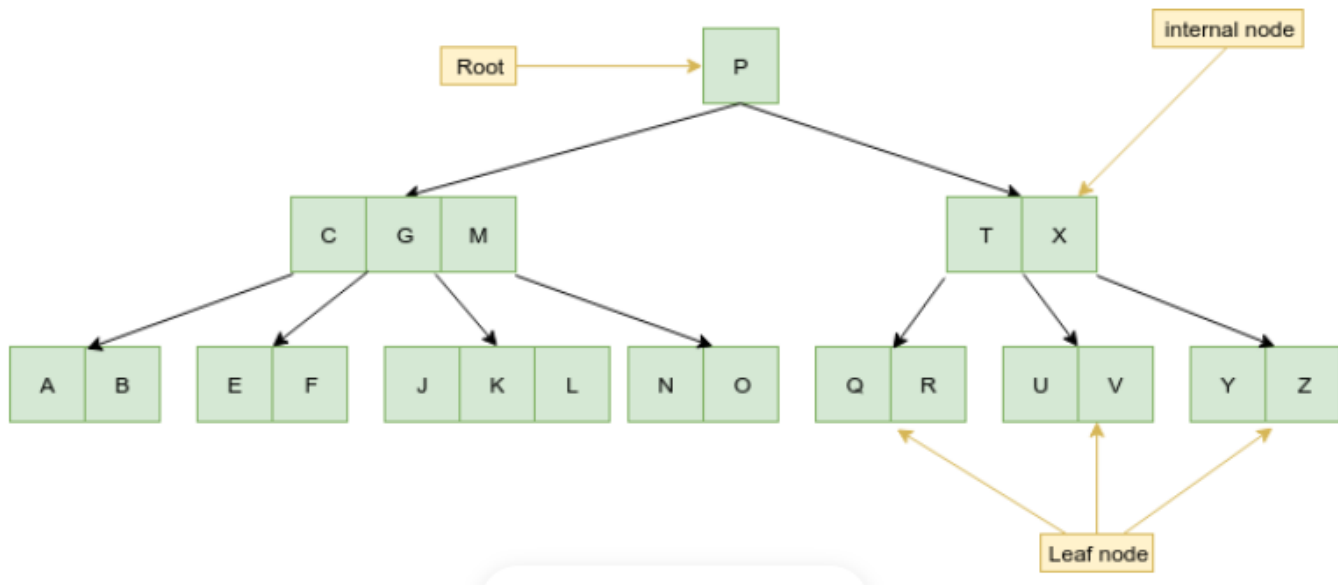
Visual Representation of the Query execution



So, indexing just cuts down the time complexity from $O(n)$ to $O(\log n)$.

More from the software perspective how the BTree works:

B-trees are often the data structure of choice to implement database indexes. A B-tree is a generalization of a binary search tree (BST). A BST is a sorted data structure composed of “nodes” (which hold data) and left and right “pointers” (which direct to other nodes within the BST). A BST has exactly two “children” nodes per node, while a B-tree can have more than two child nodes per node.



But Why only B-Tree?

The main idea of using B-Trees is to reduce the number of disk accesses. Most of the tree operations (search, insert, delete, max, min, ..etc) require $O(h)$ disk accesses where h is the height of the tree. B-tree is a fat tree. The height of B-Trees is kept low by putting the maximum possible keys in a B-Tree node. Generally, the B-Tree node size is kept equal to the disk block size. Since the height of the B-tree is low so total disk accesses for most of the operations are reduced significantly compared to balanced Binary Search Trees like AVL Tree, Red-Black Tree, ..etc.

Properties of B-Tree:

- All leaves are at the same level.
- A B-Tree is defined by the term minimum degree 't'. The value of t depends upon disk block size.
- Every node except the root must contain at least $t-1$ keys. The root may contain a minimum of 1 key.
- All nodes (including root) may contain at most $2t-1$ keys.
- The number of children of a node is equal to the number of keys in it plus 1.
- All keys of a node are sorted in increasing order. The child between two keys k_1 and k_2 contains all keys in the range from k_1 and k_2 .

- B-Tree grows and shrinks from the root which is unlike Binary Search Tree. Binary Search Trees grow downward and also shrink downward.
- Like other balanced Binary Search Trees, the time complexity to search, insert and delete is $O(\log n)$.

Advantages:

- Speed up SELECT query
- Helps to make a row unique or without duplicates(primary,unique)
- If index is set to full-text index, then we can search against large string values. for example to find a word from a sentence etc.

Disadvantages

- Indexes take additional disk space.
- indexes slow down INSERT,UPDATE and DELETE, but will speed up UPDATE if the WHERE condition has an indexed field. INSERT, UPDATE and DELETE becomes slower because on each operation the indexes must also be updated.

What is the difference between a key and an index in DBMS? Can't you use the primary key to find any piece of data?

The analogy:

Each library(database) has books (tables). Each book(table) has pages(rows). Each page(row) has a unique page number to identify it (key-value), so you could say that the *key* to finding a certain page is through its number. But if you want to find a certain page, you don't know exactly where to look for it, right? You'll have to turn page by page until you reach the page number you were looking for. Wouldn't it be easier if you had several bookmarks that tell you on which page they were placed? If you look for page 1267 and you have a bookmark on page 1266 then you go straight to that bookmark and now you only have one page to turn to get to your page. The index is like a set of bookmarks for a book. It helps you find the desired rows much faster.